

Nigeria Avian Influenza Systematic Review

2006-2026 • all phases complete • operational record current to 10 October 2026

WORKFLOW STATUS AND COMPLETION RECORD — Current operational record: 10 October 2026

Purpose

This report replaces the workflow-status snapshot of 19 September 2026, which still showed the review at its screening stage. It records the complete journey - protocol, search, register and screening, retrieval, extraction, risk of bias, synthesis, figures, the October 2026 update search, and the manuscript build and checks - together with the controls that held, the author-owned items that remain, and the source of every number quoted. The manuscript is the deliverable; the Workbench remains the operational audit record. No number below is carried over from the superseded snapshot.

1. Status at a glance

Register	Reports / studies	Manuscript	Author items	Deadline
2,157	127 / 113	PASS	6 + 1	Sun 11 Oct
Final. 1,803 title-excluded + 342 full-text candidates + 12 duplicates.	Retained reports, linked into 113 studies.	verify.py PASS; docx validation PASSED; 96-page render with Figures 2-3.	Six highlighted placeholders and the journal choice remain.	Submission-ready package; author asked for Sunday 11 October 2026.

Every phase of the review is complete and the register is final. The manuscript passes every check the project can run: 145 references (4 flagged and disclosed), 8,530 main-text words, a 290-word abstract, and a 96-page render including appendices. Nothing further can be added to the evidence base from the sources available. What remains is a short list of author-owned items, two agent-side checks, and the choice of journal.

Phase	Status	Current evidence
0. Protocol and eligibility	COMPLETE	Frozen on 16 September 2026 and re-frozen as amendments were taken; amendments dated and graded; registration position recorded honestly, not backdated.
1. Search (all sources)	COMPLETE	Six original sources 16-19 September 2026; Google Scholar update search 28 September; strategies rebuilt with the caps removed and citation chasing closed 10 October 2026; every limitation disclosed.
2. Register and screening	COMPLETE	2,157 records final; two independent AI reviewers, blind, with a third AI agent as arbiter; agreement reported per stage.
3. Retrieval	COMPLETE	Legal open routes only; 98 reports not retrieved and disclosed with the direction of effect.
4. Extraction	COMPLETE	113 studies; dual blind extraction with adjudication against the source texts.
5. Risk of bias and certainty	COMPLETE	JBI checklists plus the seven-item molecular transparency check; 14 high-risk studies; 13 studies with molecular caps; item agreement 79.3%.
6. Synthesis	COMPLETE	SWiM groupings S1-S5; four sensitivity analyses (i)-(iv); all assertions PASS at coverage 113.

Phase	Status	Current evidence
7. Figures	COMPLETE	State x epoch map and clade/wave timeline, both built by script with provenance-validated ranges.
8. October update search	COMPLETE	17,037 records retrieved by the rebuilt strategies; citation chasing closed at 88/123 seeds with 0 errors; 52-record packet screened blind; 4 studies retained abstract-only.
9. Manuscript build and checks	COMPLETE	145 references (4 flagged, disclosed); verify.py PASS; docx validation PASSED; 96-page render including appendices; Figures 2-3 embedded.

What remains

Author-owned: the six highlighted placeholders - the declarations (funding, which also appears in the abstract, competing interests, author contributions, acknowledgements and data availability) plus the Reviewer 1 line - and the choice of journal. Agent-side and pending: re-running the reporting audits against the updated manuscript, and the reviewer blind re-score. The author asked for the commit; it was made on 10 October 2026 (nothing pushed). Sections 6 to 8 list each item with its owner and its state.

DEADLINE	Sunday 11 October 2026 is the hard deadline the author set ("I need this by sunday"), read as the submission-ready manuscript package. The review work is the agent's job: only the final deliverable goes to the author.
----------	---

2. The record: the phase-by-phase journey

Dates are those of the dated run-log and handoff entries. Agreement figures describe agreement between **AI** agents, never between humans.

Phase (dates)	Record	State
0. Protocol 16 September 2026	The protocol was frozen on 16 September 2026 and re-frozen as amendments were taken. It carries the dated eligibility anchor, the reviewer model, the screening form and the graded amendments; the freeze is consistent with the registration position and no backdating is claimed. Reviewer roles and the SWiM default were fixed here.	COMPLETE
1. Search 16-20 September 2026; Scholar update 28 September	PubMed, OpenAlex/Crossref, Scopus and ScienceDirect were last run on 16 September 2026 and the Web of Science Starter API update on 19 September (the record-level export of 20 September is not recoverable). A Google Scholar update search was run through SerpApi on 28 September, after the protocol freeze, and is reported as identification by other methods. The October rebuild and its per-source audit are in section 3.	COMPLETE
2. Register and screening 16 September - 10 October	Records were de-duplicated by DOI, PMID or normalised title plus year. Screening was two-stage (title/metadata, then full text) by two large language model agents working independently and blind to each other, with a third AI agent (the coordinating session) adjudicating every disagreement against the protocol criteria; at the title stage a record retained by either reviewer proceeded to full text. Reviewer 2 was an independent Codex agent (Epicurus) from 19 to 27 September; from 27 September Reviewer 1 was an executor agent and Reviewer 2 a separate reviewer agent. Agreement: 97.0% (kappa 0.88) at title/metadata on 1,877 records; 86.6% (kappa 0.76) at full text on 216 reports; 94.3% (kappa 0.77) on 228 records for the Scholar update search; and 92.3% (kappa 0.85) on 52 records for the October update search, whose final packet numbers are the newest layer.	COMPLETE
3. Retrieval September - 9 October	Full texts were sought only through identifier resolution and legal open-access routes (Crossref, OpenAlex, Europe PMC, PubMed, Semantic Scholar, Unpaywall, institutional repositories and Google Scholar version clusters). No CAPTCHA was bypassed, no login or shadow library was used, and institutional access was not available to the review team. The September open-route re-sweep recovered a file for 18 of the 88 reports then outstanding (10 full texts, 8 conference abstracts); three files obtained after a browser auto-passed an anti-bot check were deleted. Of the 17 records advanced by the October packet, 1 full text and 6 abstracts were retrieved and 10 were not retrievable. Reports with no verified full text were recorded as not retrieved and were not assessed for eligibility: 98 in total (88 + 10), of which about 40 were judged likely relevant - disclosed as a limitation, not silently dropped.	COMPLETE
4. Extraction 30 September - 10 October	Reports were linked into studies before extraction: links proposed by one AI agent from the full texts, checked against the texts by a second, and adjudicated by the coordinating session. The extraction form was piloted blind by two extractors on the same 10 studies. Each study was then extracted independently and blind by two AI agents, with every value tied to a verbatim locator in the source text, and disagreements adjudicated against the texts by the coordinating session. Authors were not contacted for missing data. Coverage is 113 studies, including the October additions STU-113, STU-114 and STU-115 and the abstract-only set; CAN-2139 was linked as a companion report of STU-023 rather than as a new study.	COMPLETE
5. Risk of bias and certainty 3 - 10 October	Each study was appraised with the JBI checklist matching its design (Prevalence 55, Analytical Cross-sectional 5 and Case Series 30, the last added by a dated minor amendment on 3 October 2026), with a structured note instead of a checklist for 20 secondary or modelling studies, and the 37 studies with molecular data were also assessed with the seven-item molecular transparency check M1-M7 (a "No" on accession availability or phylogenetic support caps a study's molecular claims at "suggestive"). Two AI appraisers worked independently and blind, and the coordinating session adjudicated against the texts. Item-level agreement before adjudication was 868 of 1,094 items (79.3%). Scores were never summed. The October additions raised the high-risk count from 11 to 14; 13 studies carry molecular caps.	COMPLETE
6. Synthesis 7 - 10 October	Meta-analysis was not attempted: designs, units (birds, flocks, outbreaks, sequences) and assays were too heterogeneous to combine. Synthesis followed SWiM with five groupings matched to the review question - S1 temporal, S2 geographic, S3 host, S4 molecular, S5 exposure - over epochs 2006-2010, 2011-2014, 2015-2020 and 2021-2026. Four sensitivity analyses were run: (i) excluding high-risk studies; (ii) excluding grey literature report by report; (iii) as (ii) with four reports of uncertain peer-review status also treated as grey, added during synthesis before certainty rating; and (iv) restricting state counts to states with sequence or isolate evidence, added post hoc after certainty rating. After the October update search, coverage is 113 studies, all script assertions PASS, the Table 2 counts are unchanged (27/22/15/0/5) and S5 holds 110 rows / 72 independent; sensitivity (i) was updated and (ii)-(iv) are table-level unchanged.	COMPLETE
7. Figures 9 October	Figure 2 is the state x epoch map built from the synthesis tables; Figure 3 is the clade and wave timeline. Both are built by script from the synthesis outputs, with every drawn range moved into a cited provenance file and a validation table carrying input hashes and a build date. The fix round closed the 2.3.2.1c bar at its last evidenced point with an honest continuation treatment and closed 2.3.4.4b at 2026. Both figures are embedded in the manuscript and were visually inspected in the render.	COMPLETE
8. October update search 8 - 10 October	All search strategies were rebuilt and the record caps removed. 17,037 records were retrieved and de-duplicated against the register (2,270 matched; 44 likely-new, 43 in window, after the reviewer-corrected de-duplication). Citation chasing ran on 123 seed reports and closed on 10 October (88 resolved, 3 with no DOI, 0 errors): 2,195 candidates collapsed to 1,745 records new to the register, of which 11 were relevance-positive. Those joined the rebuild hits in a 52-record packet screened blind (row 2): 35 title-excluded, 17 pursued, 10 not retrievable, 3 excluded (a news item, a negative-only survey and a conceptual GIS piece) and 4 retained as abstract-only studies. The register gained CAN-2106 to CAN-2157.	COMPLETE
9. Manuscript build and checks 9 - 10 October	The target manuscript was rebuilt with all counts updated across the abstract, Methods, Results, Discussion, Appendices A/B/C and the flow diagram, with a new Appendix A3 (October update search and PRESS self-assessment) and amendment rows for 8-9 October. Checks: 145 references, 4 flagged for genuinely incomplete metadata and disclosed (nothing invented);	COMPLETE

Phase (dates)	Record	State
	verify.py PASS (145 references listed and cited, first-citation order monotonic, all 113 studies cited, no stray study ids, 6 author placeholders, no replacement characters); docx validation PASSED with Figures 2-3 embedded; render 96 pages including appendices; abstract 290 words and main text 8,530 words. Backups were taken before the figure and update-search rebuilds.	

3. Search-source audit, per source

Source status combines the original layers of 16-19 September 2026 with the October 2026 rebuild, which removed the record caps. Where a limitation could not be repaired it is kept visible here rather than quietly dropped.

Source	Layer	Records and limitations	Reporting control
PubMed / MEDLINE	Original 16 Sep 2026 (v3); rebuild 9 Oct	Original string returned 140 records. The rebuilt string (adding "Influenza in Birds"[Mesh] and related terms) returned 209, about 49% more. Abstracts were fetched for the 51 unmatched PubMed records, which found 5 newly relevant records.	Both strings and their counts are documented; the gain from the improved string is reported, not hidden.
OpenAlex	Rebuild 9-10 Oct	10,614 records (the 1,000-record cap was removed). The reviewer pass found that OpenAlex stores DOI/PMID as full URLs, so the first de-duplication produced zero DOI matches; the corrected de-duplication was rerun. The free daily budget was exhausted (HTTP 429), blocking 43 citation-chasing seeds until the 00:00 UTC reset, after which the re-chase closed at 88/123 seeds resolved with 0 errors.	Per-source totals and the rate-limit interruption are recorded; citation chasing is now complete.
Crossref	Rebuild 9 Oct	6,000 records of 16,292 returned (the source cap was reached).	The cap is stated wherever the count is used, and is not presented as coverage of the source.
Web of Science Starter API	Original 19 Sep; rebuild 9 Oct	Original 175 records across 4 pages (175 unique UIDs, HTTP 200). Rebuild 194, with the date filter applied locally because the Starter API ignores it. Starter returns no abstracts, so records are matched on title and metadata only. The corrected facet queries of 19 September returned HTTP 429 for all three facets and no corrected facet total is claimed.	Base-query counts only; facet totals never claimed as complete.
AJOL	Rebuild 9 Oct	66 rows returned, but result pages 2-4 were byte-identical to page 1, so only 20 records were captured; the remaining 46 titles would need a browser. AJOL was not one of the original run sources.	The pagination defect is stated wherever the count is used.
Google Scholar	SerpApi; supplement 17 Sep; update search 28 Sep; rebuild skipped	The 17 September supplement (140 retrieved) cannot be reconstructed and is not relied on. The 28 September update search is the operative layer: 228 records screened, 26 retained, 190 excluded, 12 duplicates removed. In the October rebuild the SerpApi quota was spent, so Scholar was skipped (8 of 10 source calls succeeded).	Reported as supplementary identification by other methods, never as a primary database layer.
Scopus and ScienceDirect	Not rebuilt (author decision not pursued)	The 16 September strings and exports are undocumented and unrecoverable, so these sources are not part of the October rebuild and no rebuilt string exists for them.	Disclosed as an unrecoverable limitation rather than repaired silently.

Limitations of the search record, kept honest

- The original strategies were not peer reviewed with the PRESS checklist. The rebuilt October strings were self-assessed against PRESS and that assessment is reported in Appendix A3 of the manuscript, but no independent PRESS reviewer was available.
- Record caps (Crossref 6,000 of 16,292; OpenAlex and the earlier WoS layer) mean the returned sets are not exhaustive for those sources; the caps and their effect are stated wherever the counts are used.
- The Starter-tier Web of Science layer returns no abstracts, so relevance assessment on those records is title-and-metadata only.
- Google Scholar cannot be reconstructed for the 17 September supplement; only the 28 September update search is relied on, and it is reported as identification by other methods.
- No prospective registration exists: the protocol is available from the corresponding author on request, and the dated 16 September freeze is not backdated or upgraded after the fact.

4. Final accounting

The register closes at report level and at study level. Both lines are reproduced from the register and from the manuscript build, and they reconcile arithmetically.

The register at closure (report level)

Register line	Number	Basis
Register, final	2,157	Canonical identities after the October integration (CAN-2106 to CAN-2157 added).
Title/metadata excluded	1,803	1,768 from the original arms + 35 from the October packet.
Full-text candidates	342	325 from the original arms + 17 from the October packet (88 + 10, 123 + 4, 114 + 3 all reconcile).
- retained (provisional full-text eligible)	127	123 from the original arms + 4 retained abstract-only in October.
- excluded at full text	117	114 from the original arms + 3 excluded in October (news item, negative-only survey, conceptual GIS piece).
- not retrieved	98	88 from the original arms + 10 from the October packet; disclosed with the direction of effect.
Duplicates removed	12	Collapsed during de-duplication; retained in the audit trail, not as evidence records.
Check	$1,803 + 342 + 12 = 2,157$	And $342 = 127 + 117 + 98$.

From reports to studies

The 127 retained reports were linked into 113 studies: five reports carry more than one study, and companion reports were attached to the study they duplicate rather than counted again (CAN-2139 was linked to STU-023 on author evidence). Every data stream reused across reports is counted once, under one primary study. For continuity with the pre-October layer: the register held 2,105 records (1,768 title-excluded, 325 full-text candidates = 123 retained / 114 excluded / 88 not retrieved, 12 duplicates); the October integration added 52 records, of which 35 were title-excluded and 17 pursued.

Manuscript numbers (verified on the 10 October build)

Quantity	Value	Verification
References	145	4 flagged and disclosed (incomplete metadata in the new records); verify.py PASS with 145 listed and all cited, first-citation order monotonic.
Abstract	290 words	Counted from the built document between the abstract and Introduction headings.
Main text	8,530 words	Introduction to Conclusions, excluding tables and captions.
Render	96 pages	Includes Appendices A-F; Figures 2-3 embedded and visually inspected.
Studies	113 of 113 cited	No uncited studies, no stray study ids, no replacement characters.
Author placeholders	6	Highlighted for the author's own words; drafts prepared but not applied.
Build checks	verify.py PASS; docx validation PASSED	Run against the submission-ready build after the October update search and the figure integration.

5. Controls that held (honesty record)

These controls constrained what the review could claim. They held to the end and each is recorded in the run log, the handoff or the improvement-plan progress log.

Control	How it held
Legal open routes only	Retrieval used identifier resolution and legal open-access routes only. No CAPTCHA was bypassed, no login or shadow library was used, and no institutional access was claimed. Retrieval closed at what legal open routes yield; reports obtainable only otherwise were recorded as not retrieved and disclosed.
One attempt per bot gate	Three reports whose files had been obtained after a browser auto-passed an anti-bot check were deleted and moved to the author-download list. From then on the rule was one attempt per bot gate, never retried and never auto-passed.
No author contact	Author contact for unretrieved reports was not pursued, and authors were not contacted for missing data. This is reported as a limitation rather than worked around.
No shadow libraries	Shadow libraries were excluded at every retrieval round, including after retrieval was re-opened.
No invented human checks	No external human check of screening was performed this cycle, and none is claimed. The author's own verification is reported as author verification, with its transcription caveat, and never as an independent human screening; all agreement statistics are labelled as AI -to- AI agreement.
Registration honestly not backdated	The dated 16 September 2026 protocol freeze stands as recorded. No prospective registration is claimed, and the absence of OSF or PROSPERO registration is stated in the manuscript.
Committed only when the author asked	The author's rule is commit only when asked. The ledger, register, appraisal, plan and orientation files stayed uncommitted through every session until the author asked for the close-out commit on 10 October 2026.

6. Remaining before submission

Item	Owner	State and next action
Author declarations - the six highlighted placeholders: funding (which also appears in the abstract), competing interests, author contributions, acknowledgements, data availability, and the Reviewer 1 line	Author	Pending the author's own words. Drafts are prepared for approval and are not applied. The Reviewer 1 line exists because the identity of Reviewer 1 before 27 September 2026 is not recorded in the repository.
Target journal choice	Author	Pending. The main text is 8,530 words, so it may need trimming to the journal's limit.
Reporting audits re-run (PRISMA 2020, PRISMA-S, SWiM) against the updated manuscript	Agent	Pending on the Sunday path: the audits were written against the pre-update manuscript and must be re-run now that counts, Appendix A3 and the flow have changed.
Reviewer blind re-score of every domain (baseline: the 9 October appraisal)	Agent	Pending. The re-scoring agent must not see the improvement plan's targets.
4 flagged references	Disclosed	The new records carry genuinely incomplete metadata; they are flagged and disclosed rather than repaired by invention. No further repair is possible from the sources available.
Commit and hand over	Author	Done on the author's request: committed 10 October 2026 (nothing pushed). The final deliverable is handed over.

7. Governance and roles

The review was executed by AI agents under the author's decisions. The table states exactly who did what, so that no AI task is presented as human work.

Role	Who	Scope as recorded
Author	M.O. - the only human decision-maker	Set the protocol and the three verification eligibility rules (secondary analyses, official outbreak records, reprinted tables); verified 429 AI screening decisions - all full-text inclusions and exclusions, all duplicate removals, a seeded 10% sample of title-stage exclusions and all adjudication overrides - changing 6 of the 107 priority decisions and later reconfirming all 107 without change; gave the final decisions at report-to-study linkage and extraction; owns the declarations, the Reviewer 1 line and the journal choice.
Reviewer 1	An AI executor agent	Primary screening, full-text assessment, extraction and appraisal decisions. From 27 September 2026 Reviewer 1 was an executor agent; the identity of Reviewer 1 before 27 September is the one open governance line and is marked for the author to confirm.
Reviewer 2	An independent AI reviewer agent	Second, blind decision at every stage. Reviewer 2 was an independent Codex agent (Epicurus) from 19 to 27 September 2026, and a separate reviewer agent thereafter (Claude Sonnet 5, and Claude Sonnet 5.5 from 29 September 2026).
Arbiter	The coordinating AI session	Adjudicated every disagreement against the protocol criteria and the source texts; ran the register, the retrievals, the synthesis reruns and the manuscript rebuilds.
Appraisal	Two independent AI appraiser agents	Scored the completed review on 9 October 2026, blind to the improvement plan and its targets, with main-session verification of every material claim.

DISCLOSURE	Every screening, full-text, extraction, risk-of-bias, appraisal and synthesis decision in this review was made by AI agents. No human performed AI work and no human screening check is claimed; the agreement statistics are AI-to-AI agreement. The single human role is the author - protocol, eligibility rules, verification of the 429 AI decisions, and final decisions.
------------	---

8. Provenance of the key numbers

Every figure printed in this report, with the project file it is taken from. Where the files differ, the newest layer (the register summary and the top of the run log) is used.

Number(s)	Source
2,157 records; 1,803 / 342 / 12	Workbench/15_RUN_LOG.md (top, "W7 done" entry); Workbench/56_SCREENING_RECONCILIATION.json (metadata.status)
127 retained / 117 excluded / 98 not retrieved	Workbench/15_RUN_LOG.md (top); Workbench/56_SCREENING_RECONCILIATION.json (metadata.status); Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (top section)
127 reports linked into 113 studies; STU-113/114/115; CAN-2139 as a companion of STU-023	Workbench/15_RUN_LOG.md (top); Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (top section)
Title agreement 97.0%, kappa 0.88, 1,877 records	manuscript/sections/03_methods.txt (Screening); Workbench/15_RUN_LOG.md (screening reconciliation entry)
Full-text agreement 86.6%, kappa 0.76, 216 reports	manuscript/sections/03_methods.txt (Screening); Workbench/60_APPRAISAL_20261009.md (verified flow counts)
Scholar update search 94.3%, kappa 0.77, 228 records (26 retained / 190 excluded / 12 duplicates)	manuscript/sections/03_methods.txt (Screening); Workbench/56_SCREENING_RECONCILIATION.json (metadata.counts)
October packet 92.3%, kappa 0.85, 52 records; 35 title-excluded / 17 pursued /	Workbench/15_RUN_LOG.md (W3 adjudication and top); Workbench/56_SCREENING_RECONCILIATION.json (metadata.w1w2_update_20261009);

Number(s)	Source
4 retained abstract-only	Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (top section)
14 high-risk studies; 13 molecular caps; item agreement 79.3% (868/1,094)	Workbench/15_RUN_LOG.md (top, RoB entry); Workbench/60_APPRAISAL_20261009.md (extraction and RoB, verified)
JB1 Prevalence 55 / cross-sectional 5 / case series 30; 20 without a JB1 tool; 37 molecular studies; sensitivity analyses (i)-(iv)	manuscript/sections/03_methods.txt (Risk of bias; Heterogeneity and sensitivity analyses)
17,037 records retrieved; 2,270 matched; 44 likely-new (43 in window)	Workbench/15_RUN_LOG.md (W1 reviewer outcome and corrections)
PubMed 140 original / 209 improved string; OpenAlex 10,614; Crossref 6,000 of 16,292; Web of Science Starter 175 original / 194 rebuild; AJOL 66 with 20 captured; Scholar quota spent	Workbench/15_RUN_LOG.md (W1 search rebuild); Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (top section)
Citation chasing: 123 seeds, 88/123 resolved, 3 with no DOI, 0 errors; 2,195 candidates; 1,745 new; 11 relevance-positive	Workbench/15_RUN_LOG.md (W2 entry; overnight re-chase of 10 October, top entry)
18 of 88 reports recovered in the September re-sweep; 3 bot-gated files deleted	Workbench/59_IMPROVEMENT_PLAN_20261008.md (progress log, 8 October)
429 A decisions verified; 6 of 107 priority decisions changed; 423 agree / 6 disagree	manuscript/sections/03_methods.txt (verification paragraph); Workbench/15_RUN_LOG.md (register verification entries); Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (9 October section)
145 references with 4 flagged; verify.py PASS; docx validation PASSED; 96-page render; Figures 2-3 embedded	Workbench/15_RUN_LOG.md (top); Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (top section); verify.py and manuscript/build_report.json (145 references, 4 flagged)
Abstract 290 words; main text 8,530 words	verify.py run against the 10 October submission-ready build (counted between the abstract and Introduction headings, and Introduction to Conclusions excluding tables and captions)
Deadline Sunday 11 October 2026	Workbench/15_RUN_LOG.md ("Deadline (D7 answered) + Sunday critical path" entry)
Controls: legal routes only; one attempt per bot gate; no author contact; no shadow libraries; no invented human checks; not backdated; commit only on the author's request	Workbench/15_RUN_LOG.md (author directive, 9 October); Workbench/57_NEXT_SESSION_HANDOFF_20260922.md (working rules); Workbench/59_IMPROVEMENT_PLAN_20261008.md (progress log); Workbench/60_APPRAISAL_20261009.md (registration ceiling)

Prepared from the current project records on 10 October 2026. This report replaces the workflow-status snapshot of 19 September 2026. It is an operational completion record, not the manuscript, and it makes no claim beyond the sources listed above; the manuscript remains the single source of truth for the review's findings.